

Assignment Project Exam Help

Search Fundamentals

<https://powcoder.com>
Radu Nicolescu
Department of Computer Science
University of Auckland

Add WeChat powcoder
12 Aug 2020

① Distributed DFS and BFS

② ClassicalDFS

③ CidonDFS §1

④ CidonDFS §2

⑤ Bellman-Ford algorithm

⑥ Maximal Independent Set

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Distributed DFS

- Despite its **sequential** appearance, DFS can work **faster** on distributed networks!

- Classical DFS** traverses each edge (**tree** and **frond**) twice, thus time complexity = message complexity = $2|E|$

- Like other distributed versions, **Cidon DFS** traverses each **tree** edge twice (needed), but avoids all or most **fronds**, thus time complexity = $2|V| - 2$, but at the possible expense of increased message complexity = $3|E|$

- Small error in Tel's text: message complexity is **not** $4|E|...$

- Cidon DFS was designed for **async** networks, thus also works for **sync** networks (even better).

Distributed BFS

- For **sync** networks, **Echo** (aka **SyncBFS**) finds a BFS spanning tree: time complexity = $2D + 1$, message complexity = $2|E|$

- Despite its **parallel** appearance, BFS is harder to implement on **async** networks and does not get all the expected benefits (bad time or bad message complexity).

- In fact, on **async** networks, BFS looks like a “milder” version of Bellman-Ford: shows similar issues, but less severe

- AsyncBFS** [Lynch]: messages = $O(D|E|)$; time—pileups = $O(D)$; time+pileups = $O(D|V|)$

- LayeredBFS** [Lynch]: messages = $O(D|V| + |E|)$; time—pileups = $O(D^2)$

- We do not further follow these issues here...

Cidon DFS

- Cidon DFS uses two types of tokens:

- One single classical **tok** token, sent on **tree** edges and, occasionally, on **fronds**

- New **vis** notification tokens, used to mark possible **fronds**

- A **vis** token can be sent:

- ahead of **tok**

- together with **tok** (can be combined)

- alone (on fronds only)

Read more: Tel §6.4, or Cidon's original paper:

Yet Another Distributed Depth-First-Search Algorithm

http://cidon.eew.technion.ac.il/files/var/448324-cidon_dfs_87.pdf

Classical vs Cidon DFS – Examples

- In all cases, there is one **single tok** token, thus the time complexity sums the delays in this token's delivery


- In classical DFS, **frond** edge $\{2, 4\}$ is sequentially **twice** traversed by the **tok** token, which adds 2 time intervals
- In Cidon DFS §1, **frond** edge $\{2, 4\}$ is **twice** traversed by **vis** tokens, in parallel with the progress of the **tok** token, which now saves 2 time intervals (while keeping messages count)
- In Cidon DFS §2, **frond** edge $\{2, 4\}$ is overlappingly traversed
 - in one direction, by the **vis** token, in **parallel** with the progress of the **tok** token,
 - in the reverse direction, by a pair of **tok+vis** tokens (**vis** is not strictly necessary)
 - this still saves 1 time interval (but increases messages count)

Classical DFS

Assignment Project Exam Help

<https://powcoder.com>

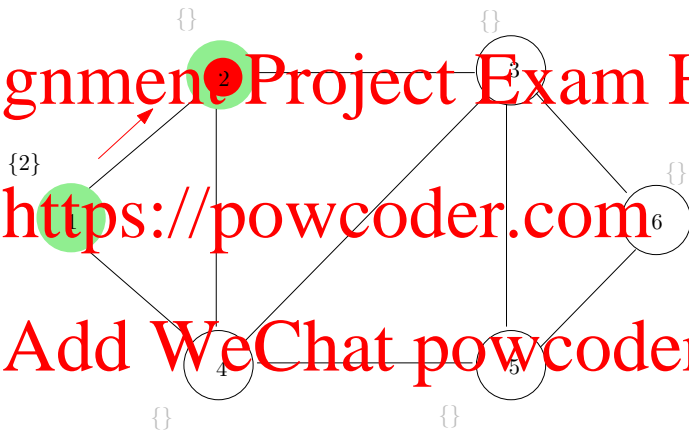
Add WeChat powcoder



Time Units = 0

Messages = 0

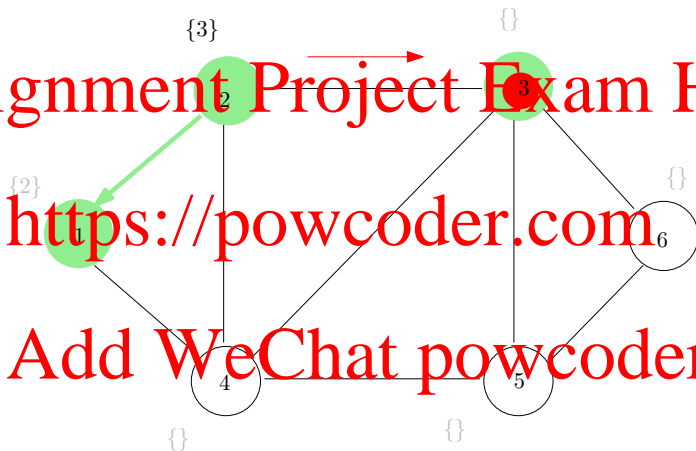
Classical DFS



Time Units = 1

Messages = 1


Classical DFS



Time Units = 2

Message = 2

Classical DFS



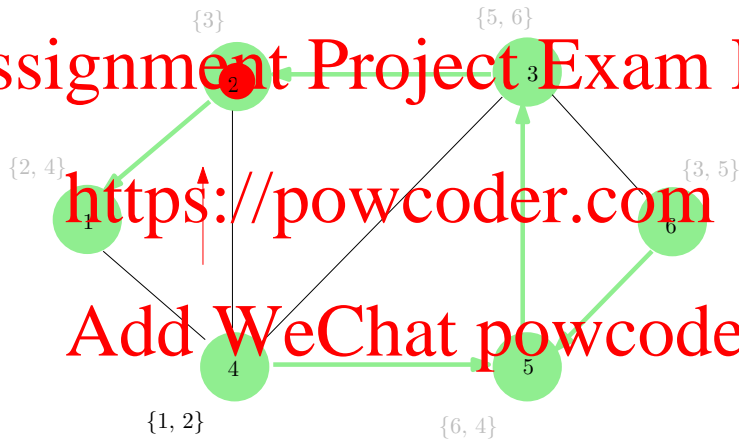
<https://powcoder.com>

Add WeChat powcoder

Time Units = 10

Message = 10

Classical DFS : 4 \rightarrow tok \rightarrow 2



Assignment Project Exam Help

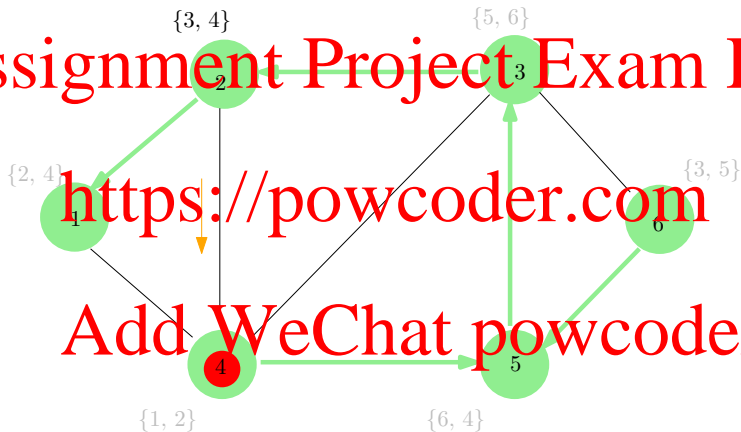
<https://powcoder.com>

Add WeChat powcoder

Time Units = 11

Message = 11


Classical DFS : $2 \rightarrow \mathbf{tok} \rightarrow 4$



Time Units = 12

Message = 12

Classical DFS




Time Units = 18 = 2M

Message = 18 = 2M

Cidon DFS §1

Assignment Project Exam Help

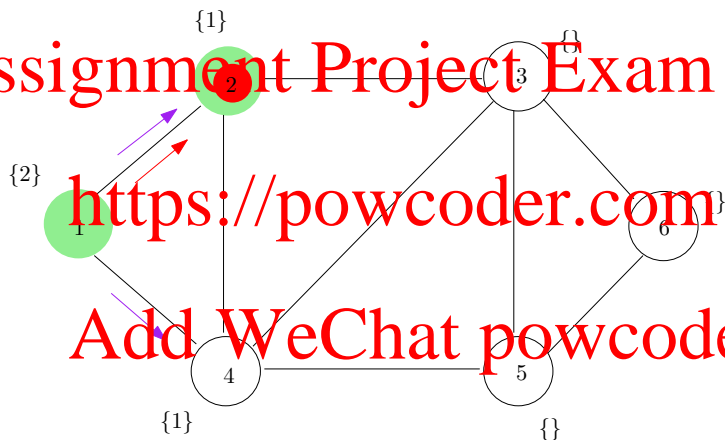


Add WeChat powcoder

Time Units = 0

Messages = 0

Cidon DFS §1



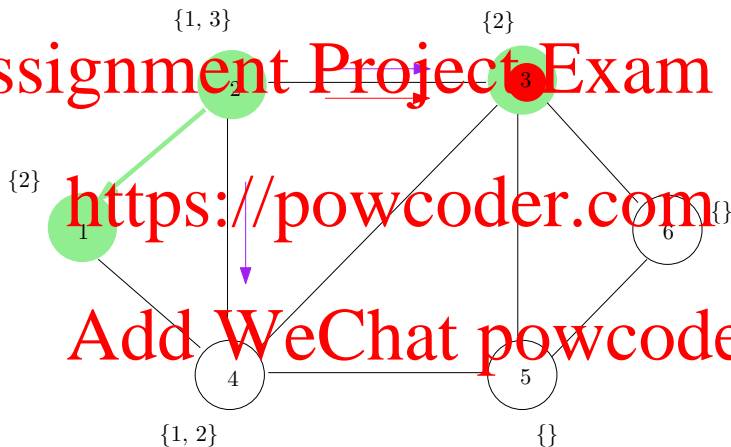
Time Units = 1

Messages = 3

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Cidon DFS §1 : $2 \rightarrow \mathbf{vis} \rightarrow 4$ 

Time Units = 2

Messages = 6


Cidon DFS §1



Time Units = 6

Messages = 16


Cidon DFS §1 : $4 \rightarrow \mathbf{vis} \rightarrow 2$



Time Units = 7

Messages = 20

Cidon DFS §1



Time Units = 10 = $2N - 2$


Messages = 23 $\leq 3M$

Cidon DFS §2

Assignment Project Exam Help

<https://powcoder.com>

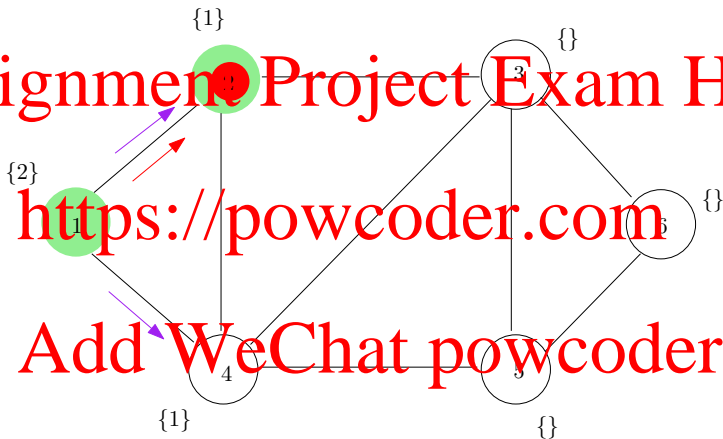
Add WeChat powcoder



Time Units = 0

Messages = 0


Cidon DFS §2



Time Units = 1

Messages = 3

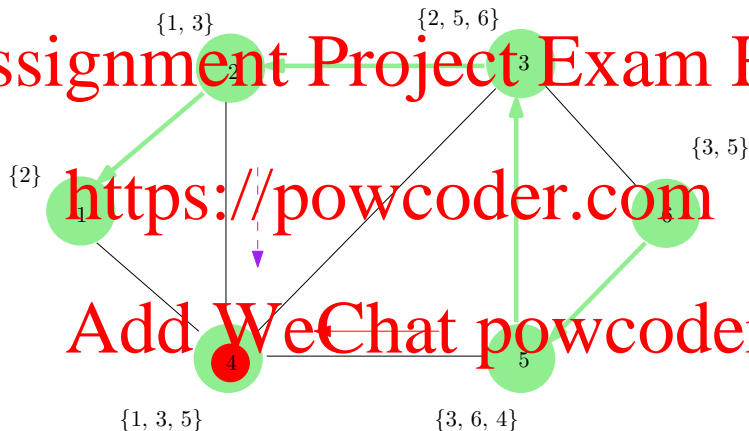
Cidon DFS §2 : $2 \dashrightarrow \mathbf{vis} \dashrightarrow 4$



Time Units = $1 + \varepsilon$

Messages = 6


Cidon DFS §2



Time Units = $1 + 5\varepsilon$

Messages = 16

Cidon DFS §2 : $2 \dashrightarrow \mathbf{vis} \dashrightarrow 4, 4 \rightarrow \mathbf{tok+vis} \rightarrow 2$




<https://powcoder.com>

Add WeChat powcoder

Time Units = $1 + 6\varepsilon$


Messages = 20

Cidon DFS §2 : $2 \rightarrow \mathbf{vis} \rightarrow 4$ 

Time Units = 2

Messages = 20

Cidon DFS §2



Time Units = $6 \leq 2N - 2$

Messages = $24 \leq 3M$

Distributed Bellman-Ford algorithm

Assignment Project Exam Help

- Classical Bellman-Ford algorithm finds all shortest paths from a single source – like Dijkstra
- Advantage for Bellman-Ford: can cope with negative weights
- Classical Dijkstra: Time complexity = $O((|E| + |V|) \log |V|)$
- Classical Bellman-Ford: Time complexity = $O(|V||E|)$
- Distributed Dijkstra: more difficult distribution...
- Distributed Bellman-Ford \approx a simple extension of Echo

<https://powcoder.com>
Add WeChat powcoder

Distributed Bellman-Ford algorithm

- Sync Bellman-Ford (“Echo ++”):

- Time complexity = $O(|V|)$
- Message complexity = $O(|V||E|)$

- <https://powcoder.com>

- Message complexity: $O(|V|^{|V|})$ (terrible worst case, but often much lower in reality)

Add WeChat powcoder
Time complexity = $O(|V|)$ – if any message is delivered in at most 1 time units – no FIFO [Tel]

- Time complexity = $O(|V|^{|V|})$ – if we consider the congestion (pileups) on FIFO channels [Lynch]
- Are these realistic? ...

Sync Bellman-Ford - Start


Problem: Find all shortest paths from node 3.

Assignment Project Exam Help

<https://powcoder.com>

- red = active
- blue = unvisited
- green = visited

Add WeChat powcoder




Sync Bellman-Ford - Round 1

Update formula: *new distance* = minimum between *old distance* and *old newly (received distance + edge length)*

<https://powcoder.com>

- Nodes 4 and 2: update distances




Sync Bellman-Ford - Round 2

Assignment Project Exam Help

<https://powcoder.com>

- Nodes 5, 6, 1:
update distances

Add WeChat powcoder




Sync Bellman-Ford - Round 3

Assignment Project Exam Help

<https://powcoder.com>

- Node 1:
update distance
(recalculation!)

Add WeChat powcoder




Sync Bellman-Ford - Round 4

Assignment Project Exam Help

<https://powcoder.com>

- Node 2:
no recalculation (but
could have been)

Add WeChat powcoder



Distributed Bellman-Ford – termination?

All nodes have successfully terminated

We can see this, but the nodes don't know this yet

How to detect and, optionally, disseminate the termination info?

- For Sync and Async Bellmann-Ford: by convergecast and, optionally, a subsequent broadcast (like Echo)
- For Sync Bellmann-Ford: by attaching a time-to-live (TTL) to the broadcast token

- Initially equal to the number of nodes $|V|$ (if this is known)
 - After receiving it, each node decrements this by 1, at each round (thus sync mode required)
 - When $TTL = 0$, each node knows that the algorithm has terminated (guaranteed)

Async Bellman-Ford – worst case (sketch, cf. Lynch §15.4)


Sketch for

- $k = 3$ (can be any number)
- $N = 2k + 2 = 8$ nodes
- $M = 3k + 1 = 10$ edges


Costs are indicated on 0 (default)

Initiator: left-most node (a)

Green: fast link; Red: slow link; Black: slow & critical if FIFO



Async Bellman-Ford – worst case (cont)



Dotted blue arrows : messages still in transit

Shapes of the shortest-so-far cost paths $\xleftrightarrow{1:1}$ base 2 numbers

Exponential message complexity $\geq 2^k = 2^{(N-2)/2} = \Omega(\sqrt{2}^N)$

Exponential time complexity if FIFO – congestion on black edge

Async Bellman-Ford – worst case

- How to explain the time complexity?

- Time complexity ~~without FIFO pileups~~ = $O(|V|)$?

- Time complexity ~~with FIFO pileups~~ = exponential?

- If the ~~last edge~~, before the rightmost node, is ~~slow enough~~, say t , then all 2^k messages can pile up in this segment

- If we consider ~~congestion~~, then these piled-up messages can be ~~successively~~ delivered at t intervals

- Total delivery time on the last edge: $2^k t = 2^{O(N)}$ time units, i.e. ~~exponential time complexity~~ [Lynch §15.4]


- This argument ~~fails~~, if we do NOT consider FIFO congestion, because then even the slowest message will not be affected by the others, and will only take a maximum 1 time unit.

Echo and Bellman-Ford – complexity highlights

- Sync Echo (aka Sync BFS) : BFS ST, no link changes, fast
- Async Echo : arbitrary ST, no link changes, but not so fast
- Sync BF : shortest paths ST, many link changes, not so fast
- Async BF : shortest paths ST, exponential link changes
 - if FIFO: worst case exponential time
- Why the worst case argument does NOT apply to Sync BF?
 - emulating slow links by extra edges exponentially increases V !

<https://powcoder.com>

Add WeChat powcoder



Maximal Independent Set

Assignment Project Exam Help

- Independent = no two neighbours

- Maximal = cannot be extended

• Maximal is ~~not~~ necessarily largest

<https://powcoder.com>

- Here we assume that nodes do **not** have IDs

- Impossible to solve with **conventional** means in an absolutely symmetric case!

- Luby's algorithm can still break the ties with **randomization** techniques

Luby's algorithm

Assignment Project Exam Help

• Sync algorithm, which works in stages, each consisting of the following 3 rounds

- 1 Each process chooses a new random value and sends it to its neighbours. Processes that have values greater than all their neighbours become **winners**

<https://powcoder.com>

- 2 **Winners** notify their neighbours. Processes that receive such messages from their neighbours become **losers**

Add WeChat powcoder

- 3 **Losers** notify their neighbours. All processes make a conceptual reshaping of the graph. Both **winners** and **losers** are **conceptually disconnected** from further participation. Remaining nodes are still **competing** and will regenerate new random values in their next stage!

Luby's algorithm

Assignment Project Exam Help

- Luby's algorithm will stop with **probability 1**, expected $O(\log n)$ rounds — proof quite hard
- The **random** values are best chosen in the range $1, 2, \dots, N^4$, will be likely distinct (but this is not necessary)
- Our diagrams will be sketched only, we won't detail all these rounds individually


<https://powcoder.com>

Add WeChat powcoder

MIS – Example

Assignment Project Exam Help

- One possible solution, formed by $\{5, 2\}$
- Another solution could be $\{5, 1, 3\}$ (which is larger)




Luby – Stage #1, Round #1

Assignment Project Exam Help

- green = hopeful
- blue = loser

- red = winner
- we only show most relevant messages



Luby – Stage #1, Rounds #2

Assignment Project Exam Help


- one winner so far: 5

- losers: 4 and 6

- losers notify their neighbours

- winners and losers are conceptually disconnected


- still competing: 1, 2 and 3



Luby – Stage #2

Assignment Project Exam Help

- a new winner: 2
- new losers, this winners' neighbours:
1 and 3
- the end



More about MIS

Assignment Project Exam Help

- Minimum Maximal Independent Set vs.
Maximum Maximal Independent Set

<https://powcoder.com>

Add WeChat powcoder

- Related readings (NOT required) – S. Butenko, PhD Thesis
<https://ufdc.ufl.edu/UFE0001011/00001/pdf>