

Assignment Project Exam Help

<https://powcoder.com>
Radu Nicolescu
Department of Computer Science
University of Auckland

Add WeChat powcoder
28 Aug 2020

① Time in Distributed Systems

Assignment Project Exam Help

② Sample Scenario

③ Lamport Logical Clock

<https://powcoder.com>

④ Vector Logical Clock

Add WeChat powcoder

⑤ Global Snapshot – Count Money

Philosophical Questions

- Does **time** exists? Is it linear? No loops?

Assignment Project Exam Help

- What is **causality**? Vs correlation.
- What is the meaning of **happened-before**? Or **concurrent**?

<https://powcoder.com>

Add WeChat powcoder



Time in distributed computing

Assignment Project Exam Help

- No global clock
- Each node has its own clock, unreliable (not totally reliable)
 - Never totally synchronised
 - Drifts, more or less
- Can we avoid these unreliable physical clocks?
- Yes, using logical clocks, that are independent of the physical time

<https://powcoder.com>

Add WeChat powcoder

Happens-before = potential causation

- Event a happens-before event b , $a \prec b$, $a \rightarrow b$, iff

① a and b occur in the same node, and a happened before b , according to the local clock

② a is a send event and b is the corresponding receive event (same message, two nodes)

- here we assume that messages are sent and received one by one, but broadcasts/multicasts can be included

③ There is an event c , s.t. $a \prec c \prec b$: transitive closure

- This happens-before ($a \prec b$) determines a partial order (creates a dag)
- For us here, happens-before \equiv potential causality

Happens-before = potential causation

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Logical time and logical clocks

Assignment Project Exam Help

- Logical time is a mapping C from events to elements of a partially ordered set, s.t. $a \prec b \Rightarrow C(a) < C(b)$

<https://powcoder.com>

- Counterfactual: $C(a) \not< C(b) \Rightarrow a \not\prec b$

- If $a \not\prec b$ then a and b are called concurrent, $a \parallel b$

Add WeChat powcoder

Logical time and logical clocks

- Logical clock is a device (algorithm) that makes these mappings in a distributed network

- Must have feature: should be determined by nodes themselves

- Must have feature: unique, i.e. one-to-one (injective)

- Good feature: same times, if “essentially” same execution

- Ideal / exact match feature: $a \prec b \Leftrightarrow C(a) < C(b)$ - NOT all clocks!

- Factual: $C(a) < C(b) \Rightarrow a \prec b$

- If the clock maps to a totally ordered set, then it is not “ideal”

Logical time in A#1?

Assignment Project Exam Help


- According to this definition, what is the logical time in A#1?
 - Message.time?
 - First four components: (.time, .code, .from, .to)?
 - What if we drop one of these?
- Is this “ideal”? Is this total?
- Is this determined by the nodes? Or by Arcs ☹
- Briefly: total logical time, but Arcs is NOT a true logical clock

<https://powcoder.com>

Add WeChat powcoder

Sample Scenario – Total orders

Two compatible totals orders for essentially the same execution



A good logical clock should generate same logical timestamps

Lamport Logical Clock – Part I

Assignment Project Exam Help

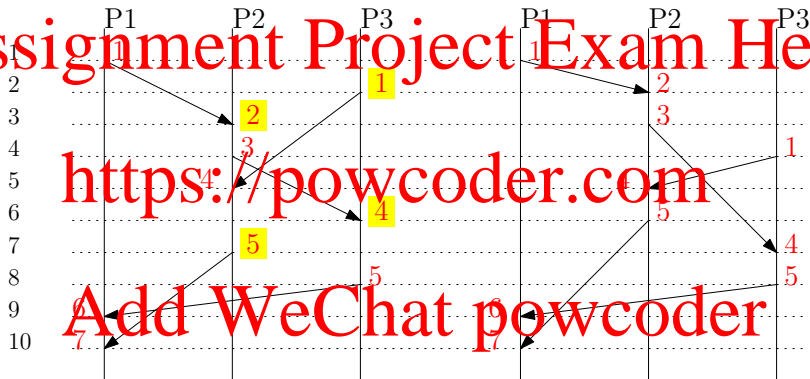
- Each processing node has its own logical clock, initialised by
 $\text{clock} = 0$
- Before each internal or send event, the clock is incremented
 $\text{clock} += 1$
- Each sent messages carries the sender's clock
- After receiving a message, the receiving node updates its clock
 $\text{clock} = \max(\text{clock}, \text{received-clock}) + 1$

<https://powcoder.com>

Add WeChat powcoder

Lamport Logical Clock – Part I

Same logical timestamps for essentially the same execution




- NOT “ideal order”: $P3 : 4 < P2 : 5$, but $P3 : 4 \not\prec P2 : 5$
- Total order, b/c timestamps have simple numerical values
- More: timestamps are NOT **unique**!

Lamport Logical Clock – Part II

Ensure **uniqueness** by adding process IDs – **lexicographic order**

Assignment Project Exam Help



- Unique, but otherwise same problems, NOT “ideal”
- NOT “ideal order”: $P3 : (1,3) < P2 : (2,2)$, but $P3 : (1,3) \not\prec P2 : (2,2)$ – in fact, these are **concurrent**

Lamport Logical Clock – Exercise (WanFok)

- Consider the following sequences of events at processes

p_1, p_2, p_3 :

```
1  $p_1 : a \ s_1 \ r_3 \ b$   
2  $p_2 : c \ r_2 \ s_3$   
3  $p_3 : r_1 \ d \ s_2 \ e$ 
```

- where s_i and r_i are corresponding send and receive events, for $i = 1, 2, 3$.

- Provide all events with Lamport's clock values.

- Answer:

```
1  $p_1 : 1 \ 2 \ 8 \ 9$   
2  $p_2 : 1 \ 6 \ 7$   
3  $p_3 : 3 \ 4 \ 5 \ 6$ 
```

Vector Logical Clock – Exact Match

- Each process keeps a vector (tuple) containing all known local logical times, e.g.

1 $V(P_2) = (v_1, v_2, v_3)$

- These tuples are ordered on each component, i.e. NOT lexicographically, e.g.

1 $(v_1, v_2, v_3) \leq (v'_1, v'_2, v'_3) \Leftrightarrow v_1 \leq v'_1, v_2 \leq v'_2, v_3 \leq v'_3,$

- The local logical time is incremented before each local or send event, e.g.

1 $V(P_2) : (v_1, v_2, v_3) \Rightarrow (v_1, v_2 + 1, v_3)$

- ...

Vector Logical Clock

Assignment Project Exam Help


- For send events, these tuples are sent along with the message

- For receive events, the new vector is the maximum of local and received tuple, with an incremented local logical time, e.g.

1
2

$$V(r_2) = (v_1, v_2, v_3), (r_1, r_2, r_3) \\ \Rightarrow (\max(v_1, r_1), \max(v_2, r_2) + 1, \max(v_3, r_3))$$

Vector Logical Clock – Complete example




(1, 1, 0) and (0, 0, 1) denote concurrent events and are incomparable!

Vector Logical Clock – Steps

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder



Vector Logical Clock – Steps

Assignment Project Exam Help




<https://powcoder.com>

Add WeChat powcoder

Vector Logical Clock – Steps


Assignment Project Exam Help



Vector Logical Clock – FIFO example / counter-example

Assignment Project Exam Help

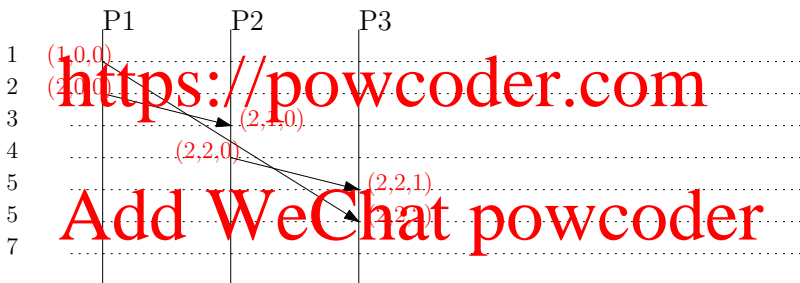
Vector logical clocks can also detect simple FIFO violations, but not Lamport



Vector Logical Clock – Causal ordering counter-example

Assignment Project Exam Help

Causal ordering ~ FIFO ++ @



How many objects are in these systems?

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder



Count Money Example – How much money in each node?

- Just after logical time 3 (3.5) or 5 (5.5)?
- Assume: initial deposits 100\$ each, arrow transfers 10\$ each.



Count Money – Algorithm

- Given Lamport time t , at each node “bank account”, sum

- The amount already there at time t

- Plus the money already sent to you, but still in transit

- I.e., for each channel
 - Keep adding all incoming money until you get the first message sent after time t – aka **marker** message

- Don't bother about money you send out after time
- Assumption #1: on each channel, the message flow is potentially infinite (it never stops, both ways)

- Assumption #2: FIFO flows ?

- Can we work around the FIFO restriction?


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Count Money Example – Amounts at Lamport times 3.5

Assignment Project Exam Help




• ? = bad, ✓ = good accounting

• local Lamport time to sum-up

	P1	P2	P3	T
	100	100	100	300
3.5 ?	90	100	90	280 ?
3.5 ✓	90 @7	110 @10	100 @10	300 ✓

Count Money Example – Amounts at Lamport times 5.5

Assignment Project Exam Help



• ? = bad, ✓ = good accounting

• local Lamport time to sum-up

	P1	P2	P3	T
	100	100	100	300
5.5 ?	90	100	90	280 ?
5.5 ✓	110 @10	100 @10	90 @10	300 ✓

Count Money Example – FIFO Discussion



<https://powcoder.com>

- FIFO fix: Add sequence numbers, for each channel! (TCP?)
- FIFO problems!
- Ensure that you sum all messages sent before the "markers"!

	P1	P2
0	100	100
1.5 ?	90	90
1.5 ???	100 @5	90 @4

	P1	P2
0	100	100
1.5 ?	90	90
1.5 ???	100 @5	90 @4
1.5 ✓	100 @5	100 @5