


SYNCHRONISERS: ATTEMPT TO SYNCHRONISE ASYNC NETWORKS (SEE LYNCH CH 16)**DO NOT COPE WITH STOPPING (OR, EVEN WORS, WITH BYZ) FAILURES****GlobSync (Global Synchronizer)**

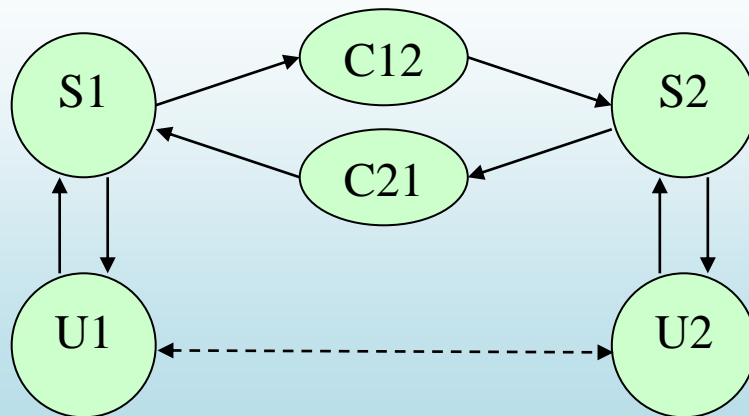
- ensures that *all* messages are sent before receiving
- collects *all* sent messages and then delivers them to their targets

**LocSync (Local Synchronizer)**


- similar to GlobSync
- but only ensures that *all neighborhood* messages are sent before receiving

SimpleSync (Simple Synchronizer)

an implementation of the LocSync



- Synchronizers need to keep the round number, explicitly or implicitly.
- Even if no faulty nodes are present, messages from different rounds may arrive interleaved (**fast**, **slow**).
- Here, process U3 may receive a round 2 message from U2, before receiving its round 1 message from U1 !



Assignment Project Exam Help

SYNCHRONISERS: CAN THEY HELP?

Add WeChat powcoder

- Using a simpler sync algorithm instead of a much more complex async !
 - E.g. Sync Distributed MST instead of the Async Distributed MST (GHS)

Assignment Project Exam Help

- Obtaining a faster runtime ?
 - E.g. could this avoid greedy choices, which may be very costly in the end ?
- What if we apply a synchroniser to Distributed Bellman-Ford ?

<https://powcoder.com>

Add WeChat powcoder

COMMUNICATION FAILURES

Assignment Project Exam Help

Is it possible to reach a distributed agreement with (unbounded) communication failures?

Add WeChat powcoder

Formal conditions.

- **Termination:**
 - All processes eventually decide.
- **Agreement:**
 - No two processes ever decide on different values.
- **Validity:**
 - If all processes start with the same initial value 1 and all messages are properly delivered, then 1 is the only one possible decision value. [WEAK]
 - **relaxation:** allow fallback on 0 if communications fail (approx. v0 could be 0)
 - and
 - If all processes start with the same initial value 0, then 0 is the only one possible decision value. [STRONG]

We also assume that a decision once taken can be immediately acted upon (by *all processes in sync*).

Diagrams

	Initial	Decision
No comm fault	$\forall 0$	$\forall 0$
Comm faults	$\forall 0$	$\forall 0$







	Initial	Decision
No comm fault	$\forall 1$	$\forall 1$
Comm faults	$\forall 1$	$\forall 1$ or $\forall 0$

	Initial	Decision
No comm fault	$\exists 0, \exists 1$	$\forall 1$ or $\forall 0$
Comm faults	$\exists 0, \exists 1$	$\forall 1$ or $\forall 0$

Fundamental result: Even under such relaxed conditions, no deterministic agreement is possible if unlimited communication failures are possible (but then, how do TCP and the internet work?)

Proof: by induction, using the given validity, agreement and termination conditions

In our diagrams, **!** denotes a no-return state, where a certain decision will inevitably follow (even if no more messages are received)



Assignment Project Exam Help

- **2 PHASE COMMIT – WEAK TERMINATION (“BLOCKING”)**
- **3 PHASE COMMIT – STRONG TERMINATION (“NON-BLOCKING”)**

Termination: All non-faulty processes eventually decide [STRONG]. (3PC)

or

If there are no failures then all processes eventually decide [WEAK]. (2PC)

Agreement: No two processes (even faulty) ever decide on different values [STRONG]. (2&3PC)

Validity: If any process starts with the initial value 0, then 0 is the only one possible decision value.

initial 0 → can start as decided 0!

and If all processes start with the same initial value 1 and there are no failures, then 1 is the only one possible decision value.

[WEAK]. (2&3PC)

Diagrams

Process	Initial	Decision
Decided, not failing	∀1	∀1 or ∀0 [†]
Decided, before failing		∀1 or ∀0 [†]
Non-decided, failed		—
Non-decided, not failing (2PC*)		—

Process	Initial	Decision
Decided, not failing	∃0	∀0
Decided, before failing		∀0
Non-decided, failed		—
Non-decided, not failing (2PC*)		—

[†] Decision 0 is allowed only if one/more processes fail

* The 2PC weak termination condition allows non-faulty process (with initial value 1) that never decide, that remain “blocked” (unless they receive some “magical” help).

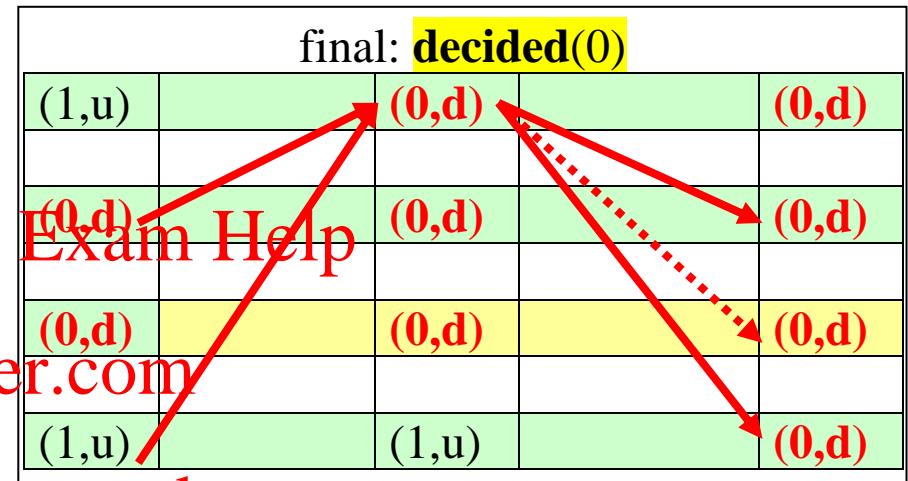
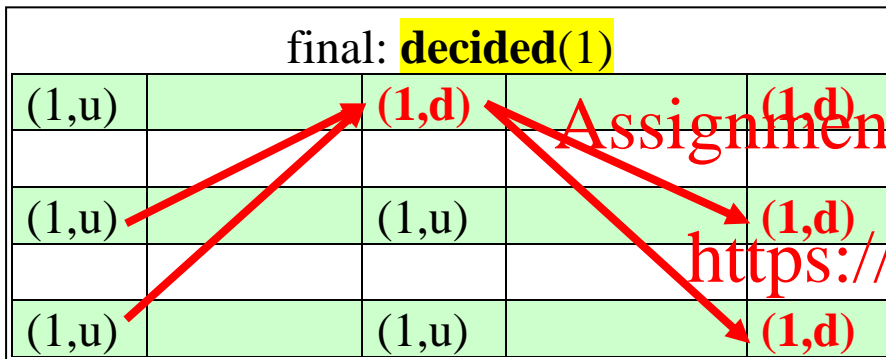
THE 2 PHASE COMMIT – WEAK TERMINATION (“BLOCKING”)

We do not discuss here:

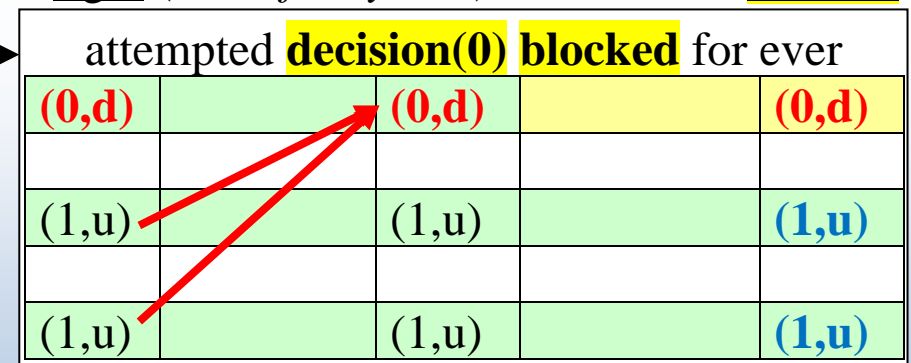
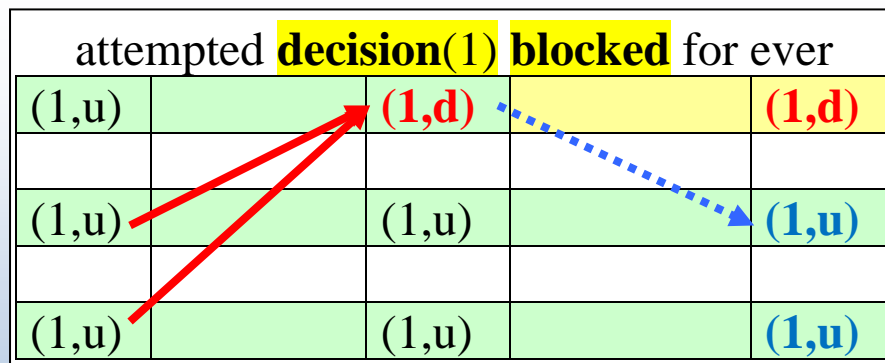
- how the processes start this
- how do they recover after failures.

- Participants: coordinator (leader), cohorts.

- States: $\text{decided}(0) \leftarrow \text{uncertain}(1) \rightarrow \text{decided}(1)$, or $+ \text{failed}$



If #1 fails, then #2&3 cannot differentiate between left and right (even if they talk), will remain **blocked**



- 2PC blocks because the leader is too greedy to decide on 1 and fails before sending its messages

THE 3 PHASE COMMIT – STRONG TERMINATION (“NON-BLOCKING”)

Assignment Project Exam Help

3PC algorithm is less greedy than 2PC and new coordinators arise in place of failed ones

States (one more, $\text{ready} \leftarrow \text{decided}(0) \leftarrow \text{uncertain}(1) \rightarrow \text{ready}(1) \rightarrow \text{decided}(1)$, or +failed)

final: **decide(1)**

(1,u)		(1,r)		(1,d)		(1,d)
(1,u)		(1,u)		(1,r)		(1,d)
(1,u)		(1,u)		(1,r)		(1,d)

final: **decide(0)**

(1,u)		(0,d)		(0,d)		
(0,d)		(0,d)		(0,d)		
(0,d)		(0,d)		(0,d)		
(1,u)		(1,u)		(0,d)		

○ State transition rules are “natural” and similar for all rounds, except:

○ In the first 3 rounds:

- Any missing message \Rightarrow **decided(0)!**
- If all are **uncertain**, they will *attempt* \Rightarrow **decided(1)!**

○ In round 4 and following:

- Any missing message is **ignored!**
- If all the rest are **uncertain**, they will *certainly* \Rightarrow **decided(0)!**

Assignment Project Exam Help

Add WeChat powcoder

Assignment Project Exam Help

(1, r) <https://powcoder.com>

Add WeChat powcoder

(0,d)	(0,d)	(0,d)								
(1,u)	(1,u)	(1,u)		(1,u)		(0,d)	(0,d)	(0,d)		
(1,u)	(1,u)	(1,u)		(1,u)		(1,u)	(0,d)	(0,d)		