

Assignment Project Exam Help

Fault Tolerant Consensus – Wrapup I

<https://powcoder.com>

Radu Nicolescu
Department of Computer Science
University of Auckland

Add WeChat powcoder

16 Oct 2020

- ① Synchronous network model

Assignment Project Exam Help

- ② Stopping failures

- ③ ElGStop

<https://powcoder.com>

- ④ Side by side

Add WeChat powcoder

- ⑤ Byzantine agreement with authentication

Synchronous network model

- All these algorithms are still based on the synchronous network model

- Even the simple stopping failure cannot be deterministically solved in the asynchronous network model

- FLP – Fischer, Lynch and Patterson

Impossibility of Asynchronous Distributed Consensus with a Single Faulty Process

- Solutions for the asynchronous model use randomisation, failure detectors (partially synchronous model)

Stopping failures model

Assignment Project Exam Help

- Much simplified version of the Byzantine agreement
- A failed process can only stop sending messages, forever (no intermittent failures, recovery not considered)
- No possibility to send confusing messages (i.e. different messages to different directions)
- The problem can be solved for any $F \leq N - 1$ (not only when $3F \leq N - 1$)

<https://powcoder.com>

Add WeChat powcoder

The Stopping agreement conditions – vs Byz

Assignment Project Exam Help

- **Termination**: all **non-faulty** processes eventually decide
- **Agreement**: no two **non-faulty** processes ever decide on different values
- **Validity**: if all **non-faulty** processes start with the **same initial value** $v \in V$, then v is the only one possible decision value
- If the processes start with **different initial values**, then the final decision could be **any of these** (as long as it is **consistent**)

<https://powcoder.com>

Add WeChat powcoder

EIGStop

- EIG tree **as in** the EIGByz, $F + 1$ messaging rounds
- recall: F can be as high as $N - 1$ not at most $(N - 1)/3$

- Top-down val()'s **as in** the EIGByz, i.e. via messaging
- No bottom-up new val() attributes

- Final decision: set **W** of all **non-null** val()'s in EIG tree
 - all values at all levels! not just leaves
 - nulls discarded, not assumed v_0


- If W is **singleton**, $W = \{v\}$, then the decision is **v**
- Otherwise, if W is **mixed**, $W = \{0, 1\}$, then the decision is **v_0**
 - no voting! no tie breaking

EIGStop example – assuming $v_0 = 1$; nulls as -

Assignment Project Exam Help

- Process #1 : init 0; decision $v_0 = 1$
- Process #2 : init 0; decision $v_0 = 1$

- Process #3 : init 1; no decision;
fails after sending one 1st round message, to #1



EIGStop example – assuming $v_0 = 1$; nulls as -

Assignment Project Exam Help


- Process #1 : init 0; decision 0

- Process #2 : init 0; decision 0

- Process #3 : init 1; no decision;
fails before sending any 1st round message

<https://powcoder.com>

Add WeChat powcoder



EIGStop example – assuming $v_0 = 1$; nulls as -

- WHAT IF scenario – NOT supported by this EIGStop protocol

Assignment Project Exam Help

What if P#3 fails before sending any 1st round out-message but would be allowed to recover and decide

- No agreement




- Process #1 : init 0; decision 0

- Process #2 : init 0; decision 0

- Process #3 : init 1; decision 1

<https://powcoder.com>

Add WeChat powcoder



OptEIGStop

Assignment Project Exam Help

- Each process sends out only **two** messages
- First: its initial choice, to all the processes
- Second: a different value, to all the processes
 - The first time it learns about a different value
 - Arbitrary choice, if there are more

<https://powcoder.com>
Add WeChat powcoder

EIGStop vs EIGByz vs 3PC – assuming $v_0 = 0$

- \times indicates a faulty process, which fails from start, before sending any 1st round message

Initial	EIGStop	EIGByz	3PC
0 0 0 0	0	0	0
0 0 0 1	0	0	0
0 0 1 1	0	0	0
0 1 1 1	0	1	0
1 1 1 1	1	1	1
0 0 0 X	0	0	0
0 0 1 X	0	0	0
0 1 1 X	0	0	0
1 1 1 X	1	1	0

EIGStop vs EIGByz vs 3PC – assuming $v_0 = 1$

- \times indicates a faulty process, which fails from start, before sending any 1st round message

Initial	EIGStop	EIGByz	3PC
0 0 0 0	0	0	0
0 0 0 1	0	0	0
0 0 1 1	1	1	0
0 1 1 1	1	1	0
1 1 1 1	1	1	1
0 0 0 X	0	0	0
0 0 1 X	1	1	0
0 1 1 X	1	1	0
1 1 1 X	1	1	0

Complexity

- EIGStop

- Rounds: $f + 1$

- Messages: $\mathcal{O}((f + 1)n^2)$ messages

- OptEIGStop messages: $\mathcal{O}(2n^2)$ messages

- EIGByz — Additional requirement: $n > 3f$

- Rounds: $f + 1$

- Messages: $\mathcal{O}((f + 1)n)$ messages

- 3PC:

- Rounds: $\mathcal{O}(f + 1)$

- Messages: $\mathcal{O}(fn)$ messages

Byzantine agreement with authentication

- Assume that each process digitally signs its messages in a total safe way, e.g. based on a reliable unbreakable PKI/DSA...

- Is this reasonable?

- Problem with certificate weaknesses: What if a powerful Byzantine faulty process is able to forge such signatures?

- Problem with authority: What if the certification authority itself is hacked or even turns into a Byzantine process?

- Anyway, assuming that such digital signatures are totally safe, Byzantine faulty nodes are not able to wreak much more havoc than a stopped process

- EIGStop can be adapted to solve the (slightly different) Byzantine agreement with authentication

- Faster/better/more general algorithms possible...